Slide Set #7

CLRM: further topics

Pietro Coretto
pcoretto@unisa.it

Econometrics
Master in Economics and Finance (MEF)
Università degli Studi di Napoli "Federico II"

Version: Monday 8th February, 2021 (h14:45)

Notes

Conditional and unconditional modeling

Suppose we want to estimate the following additive regression model

$$\boldsymbol{y} = f(\boldsymbol{X}; \boldsymbol{\beta}) + \boldsymbol{\varepsilon}$$

with the unknown parameter vector $\boldsymbol{\theta} = (\boldsymbol{\beta}, \sigma^2)$.

Two cases:

- Unconditional modeling: we treat $\boldsymbol{X}$ as random as it is the unobservable $\boldsymbol{\varepsilon}$. We do not assume knowledge about *information* generated by $\boldsymbol{X}$.
- Conditional modeling: we treat $\boldsymbol{X}$ as random, but we assume knowledge of the *information generated* by $\boldsymbol{X}$. We still do not assume knowledge of the information generated by $\boldsymbol{\varepsilon}$.

Remark: the model based on A1–A5 is a conditional model. We developed conditional inference.

Notes

In 2nd part of the course we consider large sample inference for a more general unconditional linear model. Why large samples?

Consider the OLS for the CLRM, we showed $\mathrm{E}[\boldsymbol{b} \,|\, \boldsymbol{X}] = \boldsymbol{\beta}$

Apply the "*law of total variance*" formula[1]

$$\begin{aligned}\mathrm{Var}[\boldsymbol{b}] &= \mathrm{E}[\mathrm{Var}[\boldsymbol{b} \,|\, \boldsymbol{X}]] + \mathrm{Var}[\mathrm{E}[\boldsymbol{b} \,|\, \boldsymbol{X}]] \\ &= \mathrm{E}[\mathrm{Var}[\boldsymbol{b} \,|\, \boldsymbol{X}]] + \mathrm{Var}[\boldsymbol{\beta}] \\ &= \mathrm{E}[\sigma^2 (\boldsymbol{X}'\boldsymbol{X})^{-1}] + 0 \\ &= \sigma^2 \,\mathrm{E}[(\boldsymbol{X}'\boldsymbol{X})^{-1}]\end{aligned}$$

Unconditional inference with OLS

In order to use OLS to make *unconditional inference*, we need to know $\mathrm{E}[(\boldsymbol{X}'\boldsymbol{X})^{-1}]$. We will use large sample theory to estimate these population moments using WLLN-type theorems

[1] see: https://en.wikipedia.org/wiki/Law_of_total_variance

Notes

Maximum Likelihood (ML)

Likelihood inference requires knowledge of the *joint density* of the *observable random variables* of the model.

Let $h(\boldsymbol{y}; \boldsymbol{X}; \boldsymbol{\theta})$ be the joint density of $\boldsymbol{y}$,and $\boldsymbol{X}$s, then

$$\hat{\boldsymbol{\theta}}_{\mathsf{ML}} := \arg\max_{\boldsymbol{\theta}} \; \log(h(\boldsymbol{y}; \boldsymbol{X}; \boldsymbol{\theta}))$$

Usually it is assumed that

- $\boldsymbol{y} \,|\, \boldsymbol{X}$ has joint density $f(\boldsymbol{y} \,|\, \boldsymbol{X}; \boldsymbol{\theta})$
- $\boldsymbol{X}$ has joint density $g(\boldsymbol{X})$ *not depending* on the unknown $\boldsymbol{\theta}$

This allows to factorize the joint density:

$$h(\boldsymbol{y}, \boldsymbol{X}; \boldsymbol{\theta}) = f(\boldsymbol{y} \,|\, \boldsymbol{X}; \boldsymbol{\theta}) g(\boldsymbol{X})$$

Notes

Which implies

$$\log(h(\boldsymbol{y}, \boldsymbol{X}; \boldsymbol{\theta})) = \log(f(\boldsymbol{y} \,|\, \boldsymbol{X}; \boldsymbol{\theta})) + \log(g(\boldsymbol{X}))$$

Therefore

$$\arg\max_{\boldsymbol{\theta}} \ \log(h(\boldsymbol{y}; \boldsymbol{X}; \boldsymbol{\theta})) = \arg\max_{\boldsymbol{\theta}} \ \log(f(\boldsymbol{y}, \boldsymbol{X}; \boldsymbol{\theta}))$$

Thus, ML under the unconditional (joint) model is equivalent to the one obtained under the conditional specification.

Now we will show that the connection between the OLS and ML for the CLRM model under A1–A5.

Notes

ML for the CLRM + Gaussian errors

Under (A5)

$$\boldsymbol{\varepsilon} | \boldsymbol{X} \ \sim \ \mathcal{N}(\boldsymbol{0}, \sigma^2 \boldsymbol{I}_n) \ \implies \ \boldsymbol{y} \,|\, \boldsymbol{X} \ \sim \ \mathcal{N}(\boldsymbol{X}\boldsymbol{\beta}, \sigma^2 \boldsymbol{I}_n)$$

here $f(\boldsymbol{y}, \boldsymbol{X}; \boldsymbol{\beta})$ is the multivariate normal density function with mean $\boldsymbol{X}\boldsymbol{\beta}$ and varcov matrix $\sigma^2 \boldsymbol{I}_n$

$$f(\boldsymbol{y} \,|\, \boldsymbol{X}; \boldsymbol{\theta}) = (2\pi\sigma^2)^{-\frac{n}{2}} \exp\left(-\frac{1}{2\sigma^2}(\boldsymbol{y} - \boldsymbol{X}\boldsymbol{\beta})'(\boldsymbol{y} - \boldsymbol{X}\boldsymbol{\beta})\right)$$

and

$$\begin{aligned}\log(f(\boldsymbol{y} \,|\, \boldsymbol{X}; \boldsymbol{\theta})) &= -\frac{n}{2}\log(2\pi) - \frac{n}{2}\log(\sigma^2) - \frac{1}{2\sigma^2}(\boldsymbol{y} - \boldsymbol{X}\boldsymbol{\beta})'(\boldsymbol{y} - \boldsymbol{X}\boldsymbol{\beta}) \\ &\propto -n\log(\sigma^2) - \frac{1}{\sigma^2} S(\boldsymbol{\beta})\end{aligned}$$

Notes

therefore

$$\hat{\boldsymbol{\theta}}_{\mathsf{ML}} := \arg\max_{\boldsymbol{\theta}} \left(-n \log(\sigma^2) - \frac{1}{\sigma^2} S(\boldsymbol{\beta}) \right)$$

The optimization problem is separable wrt $\boldsymbol{\beta}$ and σ^2, hence

$$\arg\max_{\boldsymbol{\beta}} \left(-n \log(\sigma^2) - \frac{1}{\sigma^2} S(\boldsymbol{\beta}) \right) = \arg\min_{\boldsymbol{\beta}} S(\boldsymbol{\beta})$$

$\implies \hat{\boldsymbol{\beta}}_{\mathsf{ML}} = \hat{\boldsymbol{\beta}}_{\mathsf{OLS}}$

It can be shown that

$$\hat{\sigma}^2_{\mathsf{ML}} = \frac{1}{n} \boldsymbol{e}'\boldsymbol{e} = \frac{1}{n} \sum_{i=1}^{n} e_i^2$$

Notes

Under CLRM + A5 (normal and spherical errors)

- for coefficients ML≡OLS,
- ML for σ^2 is biased (in finite samples)
- Cramer-Rao inequality: under regularity conditions, the ML attains the variance lower bound

Notes

Interaction terms

Consider the model

$$Y = \beta_1 + \beta_2 X_2 + \beta_3 X_3 + \ldots + \text{error}$$

Introduce the *interaction term* between X_2 and X_3

$$Y = \beta_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_2 X_3 + \ldots + \text{error}$$

Marginal effects wrt X_2 and X_3

$$\frac{\partial \, \mathrm{E}[Y \mid \boldsymbol{X}]}{\partial X_2} = \beta_2 + \beta_4 X_3 \qquad \frac{\partial \, \mathrm{E}[y \mid \boldsymbol{X}]}{\partial X_3} = \beta_3 + \beta_4 X_2$$

Observe: each marginal effect depends linearly on the level of the parent regressor in the interaction. What happens if $\beta_4 = 0$?

Notes

Suppose we observe a sample of workers that went through a training program, we want to study $\Delta\texttt{Salary}$:

$$\Delta\texttt{Salary} = \beta_1 + \beta_2 \texttt{TrainHours} + \beta_3 \texttt{SchoolYears} + \ldots + \varepsilon$$

We want to test whether the marginal effect of the train program ($\texttt{TrainHours}$) depends linearly on schooling ($\texttt{SchoolYears}$)

$$\begin{aligned}\Delta\texttt{Salary} = \beta_1 &+ \beta_2 \texttt{TrainHours} + \beta_3 \texttt{SchoolYears} \\ &+ \beta_4 \texttt{TrainHours} \times \texttt{SchoolYears} + \ldots + \varepsilon\end{aligned}$$

This is equivalent to test $\mathrm{H}_0 : \beta_4 = 0$

Notes

Dummy variables

A *dummy variable* is usually binary variable used to identify groups. Sometimes it can be a polytomous variable (more than two categories), and these introduce a finer partition of the data units.

Example: difference in means

Assume two groups (1 and 2), a response variable y can is modeled as

$$\begin{aligned} \text{Group 1:} \quad & Y_i = \mu + \varepsilon_i \\ \text{Group 2:} \quad & Y_i = \mu + \delta + \varepsilon_i \end{aligned}$$

We might be interested in testing whether $\delta = 0$.

Define a *dummy* variable

$$D_i := \begin{cases} 0 & \text{if } i \in \text{Group 1} \\ 1 & \text{if } i \in \text{Group 2} \end{cases}$$

Notes

rewrite the model as

$$Y_i = \mu + \delta D_i + \varepsilon_i \tag{7.1}$$

Assume as usual $\mathrm{E}[\varepsilon_i | D] = 0$. Observe

$$\begin{aligned} \mathrm{E}[Y_i \,|\, D_i = 0] &= \mu \\ \mathrm{E}[Y_i \,|\, D_i = 1] &= \mu + \delta \end{aligned}$$

$$\implies \delta = \mathrm{E}[D_i \,|\, d_i = 1] - \mathrm{E}[y_i \,|\, D_i = 0].$$

If we fit (7.1) and we perform the default test for $\mathrm{H}_0 : \delta = 0$, this is equivalent to test H_0 : "*mean difference across groups* $=$ *0*".

You may want to verify that the OLS estimator for the two parameters of model (7.1) are: $\hat{\mu} = \overline{y}_1$ and $\hat{\delta} = \overline{y}_2 - \overline{y}_1$. Where $\overline{y}_1$ and $\overline{y}_2$ are the sample means within the two groups.

Notes

Example: education and income

We want to model `Income`, `Age` and education. We have information on the highest level of education attained = {BA, MSc, PhD}.

We can transform the 3-levels categorical variable into 2 dummy variables $\{M, P\}$

Max school degree	Dummy	
	M_i	P_i
i has BA	0	0
i has MSc	1	0
i has PhD	0	1

The case "i has BA" is called *baseline*-group or *control*-group, less common is the terminology *reference*-group

Notes

First, we want to model that each degree provides an additional increase in income

$$\texttt{Income}_i = \beta_1 + \beta_2 \texttt{Age}_i + \beta_3 \texttt{M}_i + \beta_4 \texttt{P}_i + \text{error}_i \tag{7.2}$$

Observe: the notion of partial effects is different with non-continuous variables

- β_3: measures $\Delta\,\mathrm{E}[\texttt{Income}|\boldsymbol{X}]$ when an individual goes from BA to an MSc (w.a.e.e.)
- β_4: measures $\Delta\,\mathrm{E}[\texttt{Income}|\boldsymbol{X}]$ when an individual goes from BA to an PhD (w.a.e.e.)

Therefore the coefficient of a *dummy-coded* level expresses the variation of $\mathrm{E}[Y \mid \boldsymbol{X}]$ wrt the *baseline* level (w.a.e.e.)

Notes

Model (7.2) assumes that your school degree plays the same role for the expected income whatever it's your age.

But if you have a PhD and you are 27 this has an expected effect on your income that is different wrt to somebody else that is a 55 with a PhD.

It is reasonable to assume that a school degree affects the Income proportionally with Age. Consider

$$\texttt{Income}_i = \beta_1 + \beta_2 \texttt{Age}_i + \beta_3 \texttt{M}_i + \beta_4 \texttt{P}_i + \beta_5 \texttt{Age}_i \times \texttt{M}_i + \beta_6 \texttt{Age}_i \times \texttt{P}_i + \text{error}_i \quad (7.3)$$

$$\frac{\partial\, \mathrm{E}[\texttt{Income} \,|\, \boldsymbol{X} \;\&\; \texttt{M} = 0 \;\&\; \texttt{P} = 0]}{\partial \texttt{Age}} = \beta_2$$

$$\frac{\partial\, \mathrm{E}[\texttt{Income} \,|\, \boldsymbol{X} \;\&\; \texttt{M} = 1]}{\partial \texttt{Age}} = \beta_2 + \beta_5$$

$$\frac{\partial\, \mathrm{E}[\texttt{Income} \,|\, \boldsymbol{X} \;\&\; \texttt{P} = 1]}{\partial \texttt{Age}} = \beta_2 + \beta_6$$

Notes

Log-linear models

(aka constant elasticity models)

Consider the non-linear and non additive model

$$y = e^{\beta_1} z_2^{\beta_2} z_3^{\beta_3} e^{\varepsilon} \quad (7.4)$$

This can be *log-linearized* as follows

$$\log y = \beta_1 + \beta_2 \log z_2 + \beta_3 \log z_3 + \varepsilon \quad (7.5)$$

Observe that the partial effect computed on (7.5) corresponds to the elasiticity of the function (7.4):

$$\frac{\partial\, \mathrm{E}[\log y \,|\, z_2, z_3]}{\partial \log z_k} = \beta_k = \frac{\partial y}{\partial z_k} \frac{z_k}{y}$$

In log-linear models the coefficients are elasticities. For this reason the log-linear model is also known as the "*constant elasticity model*".

Notes

Semilog-linear models

(models for growth rates)

Consider the nonlinear model with additive errors

$$y = e^{\beta_1 + \beta_2 \texttt{Time} + \beta_3 z_3 + \varepsilon}$$

and its linearized version

$$\log y = \beta_1 + \beta_2 \texttt{Time} + \beta_3 z_3 + \varepsilon$$

Observe

$$\frac{\partial \, \mathrm{E}[\log y \,|\, z_3]}{\partial \texttt{Time}} = \beta_2$$

Autonomous growth rate

β_2 = *expected autonomous growth rate in* y. That means, β_2 is the portion of growth which is only due the evolution of the time.

Notes

Multicollinearity

Perfect multicollinearity

a regressor is an exact linear copy of another regressor, therefore the design matrix $\boldsymbol{X}$ is rank-deficient, and $\boldsymbol{X}'\boldsymbol{X}$ cannot be inverted

Multicollinearity

there regressors highly correlated each other. The effects are:

- the scatter of the regressors is highly concentrated, which means that $\det(\boldsymbol{S}_{\boldsymbol{xx}})$ is close to zero, and computation of $(\boldsymbol{X}'\boldsymbol{X})^{-1}$ may be numerically unstable
- whatever depends on $(\boldsymbol{X}'\boldsymbol{X})^{-1}$ (e.g. standard errors, default tests, etc) is not reliable

Multicollinearity is essentially a numerical problem, but in practice there is a clear intuition about it

Notes

The effect is worse if the highly correlated regressors have also both a dominant marginal effect on the Y:

$$\texttt{Earning} = \beta_1 + \beta_2 \texttt{Age} + \beta_3 \texttt{Experience} + \beta_4 \texttt{Scholing} + \ldots$$

it may be that $\mathrm{Cor}(\texttt{Age}, \texttt{Experience}) =$ large, this means that it is difficult to identify their role precisely $\implies$ *low precision = large standard errors*

Detection (some rules of thumb)

- compute solve(det(t(X)%*%X)), if it is too low (below 2×10^{16}) you may run into troubles
- compute kappa(t(X)%*%X) (= condition number of $\boldsymbol{X}'\boldsymbol{X}$), if it is larger than 50 you may run into troubles

Unfortunately absolute *save lower-bounds* don't exist!

Cures

regularization and variable transformations (not treated here)

Notes

Heteroskedasticity

In A4 we assume $\mathrm{Var}[\boldsymbol{\varepsilon} \mid \boldsymbol{X}] = \sigma^2 \boldsymbol{I}_n$:

- the diagonal of $\mathrm{Var}[\boldsymbol{\varepsilon} \mid \boldsymbol{X}]$ is constant
- $\mathrm{Var}[\varepsilon_i \mid \boldsymbol{X}] = \sigma^2$ for all i (homoskedasticity)

Under A4 the dispersion around the regression function

- cannot depend on i, the unit index
- cannot depend on the level of one or more regressors

Homoskedasticity is a strong restriction often violated in econometric models

Notes

Example: "*Espoo housing price data*"

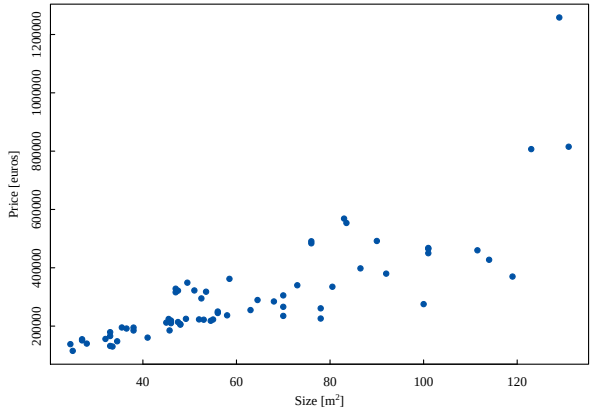

We perform OLS to estimate the following model

$$\texttt{Price}_i = \beta_1 + \beta_2 \texttt{Size}_i + \varepsilon_i$$

Notes

Estimated Regression Function

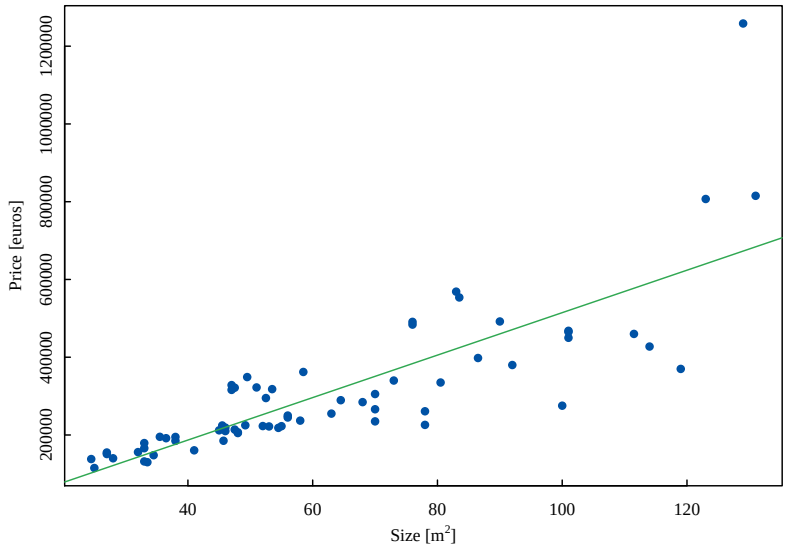

Notes

Visual Diagnostic: Residuals-vs-Fitted plot

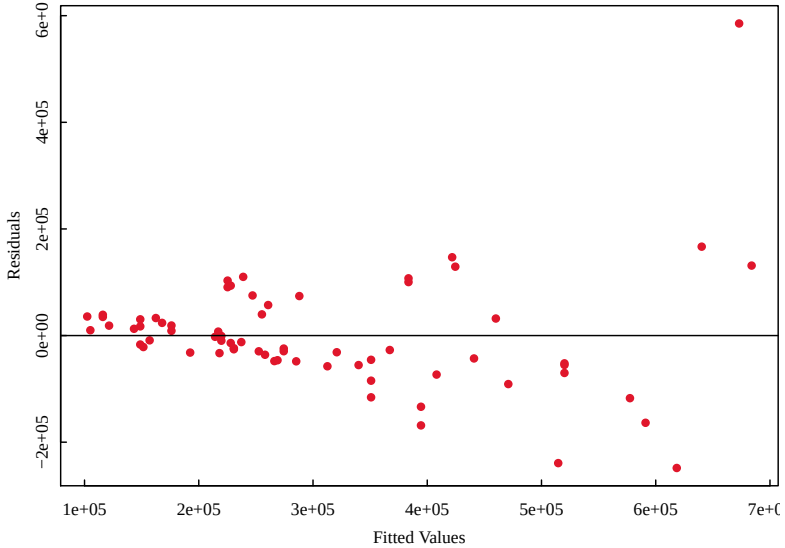

Notes

The source of the heteroskedasticity is because the variance of the error depends on the level of `Size`

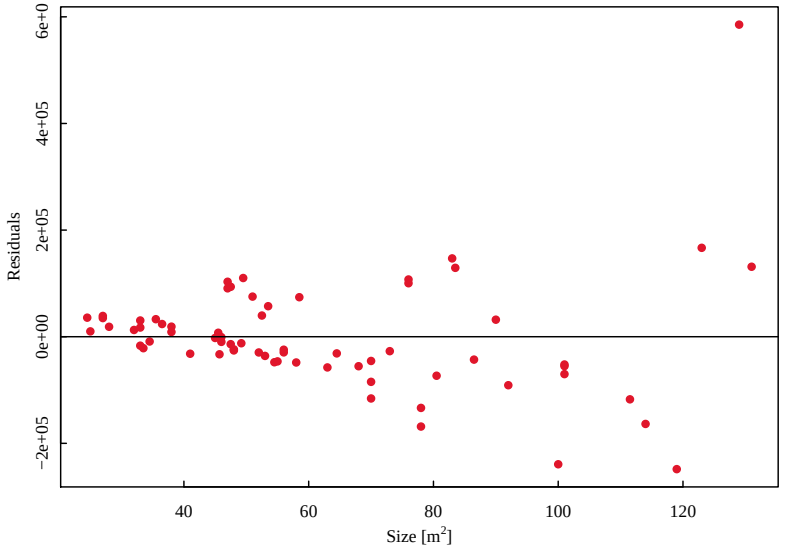

Notes

Lagged-variables and serial correlation

Lag

X_t is a RV sampled at time t, and X_{t-l} its measurement l periods before. l is called *lag*.

Suppose $X =$ `Italy's per capita GDP` measured in 2000-USD from 1970 to 2018

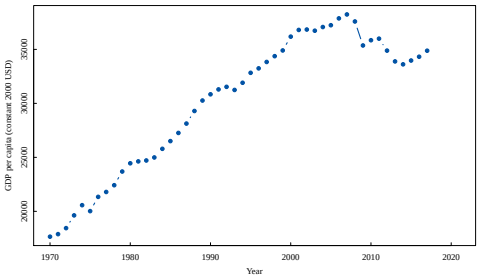

source

Notes

Imagine to construct a data matrix with three regressors X_{t-2}, X_{t-1}, X_t

$$\boldsymbol{X} = \begin{pmatrix} \texttt{GDP}_{1970} & \texttt{GDP}_{1971} & \texttt{GDP}_{1972} \\ \texttt{GDP}_{1971} & \texttt{GDP}_{1972} & \texttt{GDP}_{1973} \\ \vdots & \vdots & \vdots \\ \texttt{GDP}_{2017} & \texttt{GDP}_{2018} & \texttt{GDP}_{2019} \end{pmatrix}$$

Serial correlation (aka autocorrelation) at lag l

$$\rho(l) = \mathrm{Cor}[X_t, X_{t-l}]$$

For the GDP example, $\rho(2)$ is the correlation between the first and the third column of $\boldsymbol{X}$

Non-zero serial correlations may be be produced by *autoregressive linear models*

$$X_t = \theta_0 + \theta_1 X_{t-1} + \theta_1 X_{t-2} + \ldots + \text{error term}$$

i.e. regressors are lagged versions of the dependent variable

Notes

pairplot from the GDP data matrix $\boldsymbol{X}$

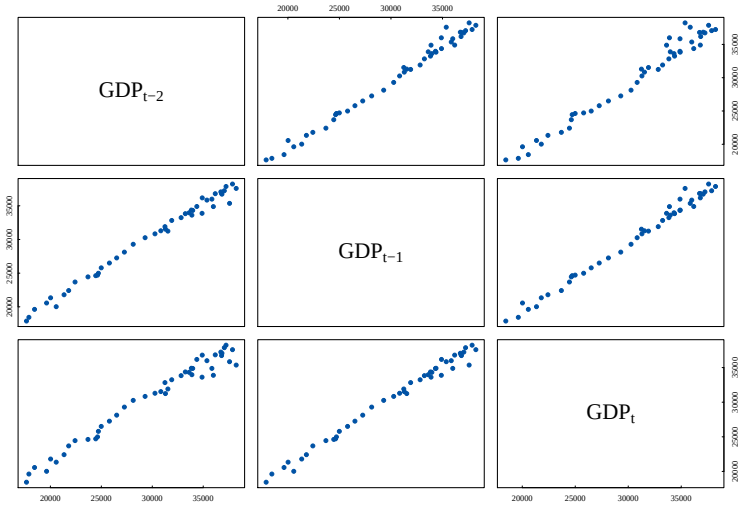

Notes

Autocorrelated errors

In A4 we assume $\mathrm{Var}[\boldsymbol{\varepsilon} \mid \boldsymbol{X}] = \sigma^2 \boldsymbol{I}_n$:

- off-diagonal elements of $\mathrm{Var}[\boldsymbol{\varepsilon} \mid \boldsymbol{X}]$ are all zero
- $\mathrm{Cov}[\varepsilon_i, \varepsilon_j] = 0$ for all $i \neq j$

Example

- $i = t = \texttt{time}$, regressors are measured as time series
- there are one or more lagged regressors

Consider the *demand model*

$$\texttt{Quantity}_t = \beta_1 + \beta_2 \texttt{Price}_t + \beta_2 \texttt{Income}_{t-1} + \ldots + \varepsilon_t \qquad (7.6)$$

Under (A4) $\mathrm{Cor}[\varepsilon_t, \varepsilon_{t-l}] = 0$ for all t and l.

(A4) implies absence of serial correlation in the error. This is often violated when we have time series sampling

Notes

Example: "*Demand for Ice Cream*" (Hildreth and Lu, 1960)

weekly observations from 18 March 1951 to 11 July 1953 on

C_t consumption of ice cream per head (in pints)
I_t average family income per week (in USD)
P_t price of ice cream (USD/pint)
X_t average temperature (in Fahrenheit)

Consider the "*demand of ice cream*" model

$$C_t = \beta_0 + \beta_1 I_t + \beta_2 P_t + \beta_3 X_t + \varepsilon_t \tag{7.7}$$

The regression function gives the expected aggregate "expenditure" for ice cream given the current information. ε_t formalizes the deviation from the expectation

$$\varepsilon_t = C_t - \mathrm{E}[C_t \,|\, \text{Information at time } t]$$

Notes

However, Hildreth and Lu (1960) found that the estimated residuals of (7.7) were correlated at lag 1. What's going on?

The aggregate expenditure measured by C_t reflects a "*rational*" learning process. The decision about C_t is likely to depend on the "*mistake*" ε_{t-1} that I did in the previous period

Suppose that this happens so that ε_t correlates with ε_{t-1}:

$$\varepsilon_t = \phi_0 + \phi_1 \varepsilon_{t-1} + u_t \tag{7.8}$$

It's easy to see that *serial correlation* in the error term implies

- $\mathrm{Cov}[\varepsilon_t, \varepsilon_{t-1}] \neq 0$ for all $t > 1$
- $\mathrm{Var}[\boldsymbol{\varepsilon} \,|\, \boldsymbol{X}]$ is not diagonal anymore
- $\mathrm{Var}[\boldsymbol{b} \,|\, \boldsymbol{X}]$ computed based on (A4) is probably wrong!

Notes

OLS with non-spherical errors

Assumption (A4-NSE): non-spherical error variance

Let $\boldsymbol{\Omega} \in \mathbb{R}^{n \times n}$ be a PSD and symmetric matrix, and

$$\mathrm{E}[\boldsymbol{\varepsilon}\boldsymbol{\varepsilon}' \mid \boldsymbol{X}] = \boldsymbol{\Omega}$$

A2 + A4-NSE $\implies$ $\mathrm{Var}[\boldsymbol{\varepsilon} \mid \boldsymbol{X}] = \boldsymbol{\Omega}$ is non-spherical.

Under A1, A2, A3, and A4-NSE, the OLS estimator $\boldsymbol{b} = (\boldsymbol{X}'\boldsymbol{X})^{-1}\boldsymbol{X}'\boldsymbol{y}$ is still unbiased (homework), but

$$\mathrm{Var}[\boldsymbol{b} \mid \boldsymbol{X}] = (\boldsymbol{X}'\boldsymbol{X})^{-1}\boldsymbol{X}'\boldsymbol{\Omega}\boldsymbol{X}(\boldsymbol{X}'\boldsymbol{X})^{-1} \tag{7.9}$$

Hence, if A4-NSE holds, all the inference methods (standard errors, tests' decision rules, etc) based on previous old formula are wrong

Notes

Dealing with non-spherical errors

In the old days there were a number of techniques to test, or correct for specific forms of heteroskedasticity or serial correlation. Today there are couple of options

Finite sample methods

- Assume that $\boldsymbol{\Omega} = g(\boldsymbol{X})$ is known, use the Generalized Least Squares Estimator (GLS) (homework)
- Under A4-NSE + Gaussian errors, the inference based on GLS will be very similar to what we have seen for OLS

Problem: $\boldsymbol{\Omega}$ is not known, and normality often fails!

Large sample methods

- consider a setup that accounts for non-spherical errors
- Gaussianity is not needed because of large sample approximations

Problem: sometimes n is not sufficiently large

Notes